

Announcements

- ◆ Homework-1 is out Tuesday, please start early
 - Use the latest Spark and Python
 - `/home/local/spark/latest/bin/spark-submit`
 - `export PYSPARK_PYTHON=python3.6`
- ◆ One additional new TA
 - Yang Zhen, zhen528@usc.edu
 - Chang Liu, liu599@usc.edu
 - Anshul Gupta, anshulg@usc.edu
 - Jing Ouyang, jingo@usc.edu
 - Xinyuan Zhou, xzhou911@usc.edu
 - Jie Ji, jiej@usc.edu

Finding Frequent Itemsets (Chapter 6)

Professor Wei-Min Shen

University of Southern California

Thanks for source slides and material to:
J. Leskovec, A. Rajaraman, J. Ullman: Mining of Massive Datasets
<http://www.mmds.org>

Frequent Itemsets and Association Rules

- ◆ Family of techniques for characterizing data: **discovery of frequent itemsets**
 - e.g., identify sets of items that are frequently purchased together

Outline:

- ◆ Introduce market-basket model of data
- ◆ Define frequent itemsets
- ◆ Discover association rules
 - Confidence and interest of rules
- ◆ A-Priori Algorithm and variations

THE MARKET-BASKET MODEL

Association Rule Discovery

Supermarket shelf management – Market-basket model:

- ◆ **Goal:** Identify items that are bought together by sufficiently many customers
- ◆ **Approach:** Process the sales data to find dependencies among items
 - Brick and mortar stores: data collected with barcode scanners
 - Online retailers: transaction records for sales
- ◆ **A classic rule:**
 - If someone buys diaper and milk, then he/she is likely to buy beer
 - Don't be surprised if you find six-packs next to diapers!

The Market-Basket Model

- ◆ A large set of **items**

- e.g., things sold in a supermarket

- ◆ A large set of **baskets**

- ◆ Each basket is a **small subset of items**

- e.g., the things one customer buys on one day

- ◆ **Want to discover Association Rules**

- People who bought $\{x,y,z\}$ tend to buy $\{v,w\}$
 - **Brick and mortar stores:** Influences setting of prices, what to put on sale when, product placement on store shelves
 - **Recommender systems:** Amazon, Netflix, etc.

Input:

<i>TID</i>	<i>Items</i>
1	Bread, Coke, Milk
2	Beer, Bread
3	Beer, Coke, Diaper, Milk
4	Beer, Bread, Diaper, Milk
5	Coke, Diaper, Milk

Output:

Rules Discovered:

$\{\text{Milk}\} \rightarrow \{\text{Coke}\}$

$\{\text{Diaper, Milk}\} \rightarrow \{\text{Beer}\}$

Market-Baskets

- ◆ Really a **general many-many mapping** (association) between two kinds of things: items and baskets
 - But we ask about connections among “items,” not “baskets.”
- ◆ The technology focuses on **common events**, **not rare events**
 - Don't need to focus on identifying ***all*** association rules
 - Want to focus on **common events**, focus pricing strategies or product recommendations on those items or association rules

Market Basket Applications (1): Identify items bought together

- ◆ **Items** = products
- ◆ **Baskets** = sets of products someone bought in one trip to the store
- ◆ **Real market baskets:** Stores (Walmart, Target, Ralphs, etc.) keep terabytes of data about what items customers buy together
 - Tells how typical customers navigate stores
 - Lets them position tempting items
 - Suggests tie-in “tricks”, e.g., run sale on diapers and raise the price of beer
 - **Need the rule to occur frequently, or no profits!**
- ◆ **Amazon’s people who bought *X* also bought *Y***
 - Recommendation Systems

Market Basket Applications (2): Plagiarism detection

◆ Baskets

➤ = Sentences?

➤ = Documents containing those sentences?

◆ Items

➤ = Sentences?

➤ = Documents containing those sentences?

Market Basket Applications (2): Plagiarism detection

- ◆ **Baskets** = sentences
- ◆ **items** = documents containing those sentences
 - Item/document is “in” a basket if sentence is in the document
 - May seem backward, but relationship between baskets and items is many-to-many
- ◆ Look for items that appear together in several baskets
 - Multiple documents share sentence(s)
- ◆ **Items (documents) that appear together too often could represent plagiarism.**

Market Basket Applications (3):

Identify related “concepts” in web documents

- ◆ **Baskets** = words? Web pages?
- ◆ **items** = words? Web pages?

Market Basket Applications (3):

Identify related “concepts” in web documents

- ◆ **Baskets** = Web pages
- ◆ **items** = words
- ◆ Baskets/documents contain items/words in the document
- ◆ Look for sets of words (items) that appear together in many documents (baskets)
- ◆ Ignore most common words
- ◆ Unusual words appearing together in a large number of documents, e.g., “World” and “Cup,” may indicate an interesting relationship or joint concept

Market Basket Applications (4): Drug interactions

- ◆ **Baskets** = patients
- ◆ **items** = drugs and side effects
- ◆ Has been used to **detect combinations of drugs that result in particular side-effects**
- ◆ **But requires extension:** Absence of an item needs to be observed as well as presence!!

Scale of the Problem

- ◆ WalMart sells 100,000 items and can store billions of baskets.
- ◆ The Web has billions of words and many billions of pages.

DEFINE FREQUENT ITEMSETS

“Support” and “Frequent Itemsets”

- ◆ **Simplest question: Find sets of items that appear “frequently” in the baskets**
- ◆ **Support for itemset I** = the number of baskets containing all items in I
 - Sometimes given as a percentage
- ◆ **Given a support threshold s , sets of items that appear in at least s baskets are called “*Frequent Itemsets*”**

Example: Frequent Itemsets

◆ Items = {milk, coke, pepsi, beer, juice}.

◆ **Support = 3 baskets.**

$B_1 = \{m, c, b\}$

$B_2 = \{m, p, j\}$

$B_3 = \{m, b\}$

$B_4 = \{c, j\}$

$B_5 = \{m, p, b\}$

$B_6 = \{m, c, b, j\}$

$B_7 = \{c, b, j\}$

$B_8 = \{b, c\}$

◆ Frequent itemsets of size 1: {m}, {c}, {b}, {j}

$\{m, b\}, \{b, c\}, \{c, j\}$.

ASSOCIATION RULES

“Association Rules” and “Confidence”

- ◆ If-then rules about the contents of baskets
- ◆ Basket I contains $\{i_1, i_2, \dots, i_k\}$
- ◆ Rule $\{i_1, i_2, \dots, i_k\} \rightarrow j$ means: “if a basket contains all of i_1, \dots, i_k then it is *likely* to contain j .”
- ◆ **Confidence** of this association rule is the probability of j given i_1, \dots, i_k
 - Ratio of support for $I \cup \{j\}$ with support for I
$$\frac{\text{support for } I \cup \{j\}}{\text{support for } I}$$
 - Support for I : number of baskets containing I

Example: Confidence

- | | |
|-----------------------|--------------------------|
| + $B_1 = \{m, c, b\}$ | $B_2 = \{m, p, j\}$ |
| - $B_3 = \{m, b\}$ | $B_4 = \{c, j\}$ |
| - $B_5 = \{m, p, b\}$ | + $B_6 = \{m, c, b, j\}$ |
| $B_7 = \{c, b, j\}$ | $B_8 = \{b, c\}$ |

◆ An association rule: $\{m, b\} \rightarrow c$

- Confidence: Ratio of support for $I \cup \{j\}$ with support for I
- Ratio of support for $\{m, b\} \cup \{c\}$ to support for $\{m, b\}$
- Confidence = $2/4 = 50\%$

➤ Want to identify association rules with high confidence

Interesting Association Rules

◆ Not all high-confidence rules are interesting

- The rule $X \rightarrow \textit{milk}$ may have high confidence for many itemsets X
 - because milk is just purchased very often (independent of X)

◆ Interest of an association rule $I \rightarrow j$: difference between its confidence and the fraction of baskets that contain j

$$\text{Interest}(I \rightarrow j) = \text{conf}(I \rightarrow j) - \text{Pr}[j]$$

- Interesting rules are those with high positive or negative interest values (usually above 0.5)
- High **positive**/**negative** interest means presence of I **encourages** or **discourages** presence of j
- Example: {coke} \rightarrow pepsi should have high negative interest

Example: Confidence and Interest

$$B_1 = \{m, c, b\}$$

$$B_2 = \{m, p, j\}$$

$$B_3 = \{m, b\}$$

$$B_4 = \{c, j\}$$

$$B_5 = \{m, p, b\}$$

$$B_6 = \{m, c, b, j\}$$

$$B_7 = \{c, b, j\}$$

$$B_8 = \{b, c\}$$

◆ Association rule: $\{m, b\} \rightarrow c$

- Confidence: Ratio of support for $I \cup \{j\}$ with support for I
- **Confidence** = $2/4 = 0.5$
- Interest: $\text{Interest}(I \rightarrow j) = \text{conf}(I \rightarrow j) - \text{Pr}[j]$
- **Difference between its confidence and the fraction of baskets that contain j**
- **Interest** = $|0.5 - 5/8| = 1/8$
 - Item c appears in $5/8$ of the baskets
 - Rule is not very interesting!

Finding Useful Association Rules

- ◆ **Question:** “find all association rules with support $\geq s$ and confidence $\geq c$ ”
- ◆ **Hard part:** finding the frequent itemsets
 - **Note:** if $\{i_1, i_2, \dots, i_k\} \rightarrow j$ has high support and confidence, then both $\{i_1, i_2, \dots, i_k\}$ and $\{i_1, i_2, \dots, i_k, j\}$ will be “frequent”
- ◆ **Assume:** not too many frequent itemsets or candidates for high support, high confidence association rules
 - Not so many that they can't be acted upon
 - Adjust support threshold to avoid too many frequent itemsets

Example: Find Association Rules with support $\geq s$ and confidence $\geq c$

$B_1 = \{m, c, b\}$

$B_2 = \{m, p, j\}$

$B_3 = \{m, c, b, n\}$

$B_4 = \{c, j\}$

$B_5 = \{m, p, b\}$

$B_6 = \{m, c, b, j\}$

$B_7 = \{c, b, j\}$

$B_8 = \{b, c\}$

◆ Support threshold $s = 3$, confidence $c = 0.75$

◆ 1) Frequent itemsets:

➤ $\{b\} \{c\} \{j\} \{m\} \{b,m\} \{b,c\} \{c,m\} \{c,j\} \{m,c,b\}$

◆ 2) Generate rules:

➤ ~~$b \rightarrow m: c=4/6$~~ $b \rightarrow c: c=5/6$ ~~$b,c \rightarrow m: c=3/5$~~

➤ $m \rightarrow b: c=4/5$... $b,m \rightarrow c: c=3/4$

➤ ~~$b \rightarrow c,m: c=3/6$~~

$$\text{conf}(I \rightarrow j) = \frac{\text{support}(I \cup j)}{\text{support}(I)}$$

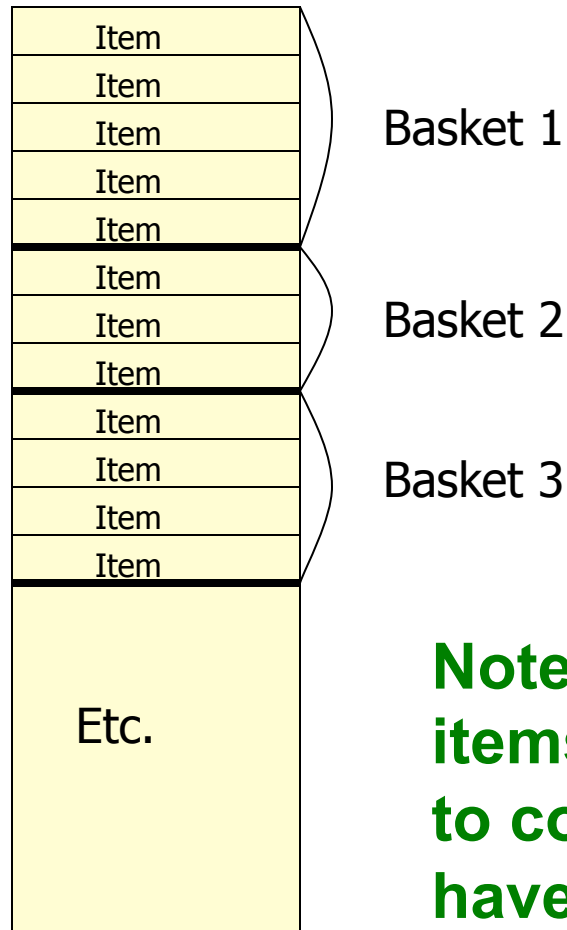
**Difficult
part is
identifying
frequent
itemsets:
algorithms
to find
them are
the focus
of this
chapter**

FIND FREQUENT ITEMSETS

Computation Model

- ◆ Typically, market basket data are kept in **flat files** rather than in a database system
 - Stored **on disk because they are very large files**
 - Stored **basket-by-basket**
 - **Goal: Expand baskets into pairs, triples, etc. as you read baskets**
 - Use k nested loops to generate all sets of size k

File Organization



Example: items are positive integers, and boundaries between baskets are -1

Note: We want to find frequent itemsets. To find them, we have to count them. To count them, we have to generate them.

Computation Model – (2)

- ◆ The true cost of mining disk-resident data is usually the **number of disk I/O's**
- ◆ In practice, association-rule algorithms read the data in *passes* – all baskets read in turn
- ◆ Thus, we measure the cost by the **number of passes** an algorithm takes

Main-Memory Bottleneck

- ◆ **For many frequent-itemset algorithms, main memory is the critical resource**
 - As we read baskets, **we need to count something, e.g., occurrences of pairs**
 - **The number of different things we can count is limited by main memory**
 - Swapping counts in/out is a disaster
 - **Algorithms are designed so that counts can fit into main memory**

Finding Frequent Pairs

- ◆ **The hardest problem often turns out to be finding the frequent pairs**
 - **Why?** Often frequent pairs are common, frequent triples are rare
 - **Why?** Probability of being frequent drops exponentially with size; number of sets grows more slowly with size
- ◆ **We'll concentrate on pairs, then extend to larger itemsets**

Naïve Algorithm

- ◆ Read file once, counting in main memory the occurrences of each pair

- Number of pairs in a basket of n items: n choose 2

$$\binom{n}{k} = \frac{n!}{k!(n-k)!}$$

- From each basket of n items, generate its $n*(n-1)/2$ pairs using two nested loops, add to the count for each pair

- First basket: (a,b), (a,c), (a,y), (b,c), (b,y), (c,y)

- Second basket: (a,b), (a,x), (a,y), (a,z), (b,x), (b,y), (b,z), ...

- Total possible number of pairs in all baskets:
 $(\#items)(\#items - 1)/2$

- ◆ Fails if $(\#items)^2$ exceeds main memory

- Remember: #items can be 100K (Wal-Mart) or 10B (Web pages)

Example: Counting Pairs

- ◆ Suppose 10^5 items
- ◆ Suppose counts are 4-byte integers
- ◆ Number of pairs of items: $10^5(10^5-1)/2 = 5*10^9$ (approximately)
- ◆ Therefore, $2*10^{10}$ (20 gigabytes) of main memory needed

Details of Main-Memory Counting

◆ Two approaches:


1. Count all pairs, using a **triangular matrix**
2. Keep a **table of triples** $[i, j, c]$ = “the count of the pair of items $\{i, j\}$ is c ”

(1) requires only 4 bytes/pair, but requires a count for each pair


Note: assume integers are 4 bytes

(2) requires 12 bytes, but only for those pairs with count > 0

Plus some additional overhead for a hashtable




Method (1):
It is a long list of "c"



Method (2)
It is a long list of
● = (i,j,c)

Triangular Matrix: (i,j) is index, c is count



Triangular-Matrix Approach – (1)

- ◆ n = total number of items
- ◆ Order each pair of items $\{i, j\}$ so that $i < j$
- ◆ Keep pair counts in lexicographic order:
 - $\{1,2\}, \{1,3\}, \dots, \{1,n\}, \{2,3\}, \{2,4\}, \dots, \{2,n\}, \{3,4\}, \dots$
- ◆ Pair $\{i, j\}$ is at position $(i-1)(n-i/2) + j - i$
 - *Every time you see a pair $\{i,j\}$ from a basket, increment the count at the corresponding position in triangular matrix*
- ◆ Total number of pairs $n(n-1)/2$; total bytes = $2n^2$
- ◆ **Triangular Matrix** requires 4 bytes (1 integer) per pair

Comparing the two approaches

◆ Approach 1: Triangular Matrix

- n = total number items
- Count pair of items $\{i, j\}$ only if $i < j$
- Keep pair counts in lexicographic order:
 - $\{1,2\}, \{1,3\}, \dots, \{1,n\}, \{2,3\}, \{2,4\}, \dots, \{2,n\}, \{3,4\}, \dots$
- Pair $\{i, j\}$ is at position $(i-1)(n-i/2) + j-i$
- Total number of pairs $n(n-1)/2$; total bytes = $2n^2$
- **Triangular Matrix** requires 4 bytes (1 integer for c) per pair

◆ Approach 2 uses 12 bytes (i, j, c) per occurring pair (but only for pairs with count > 0)

- Beats Approach 1 if fewer than 1/3 of possible pairs actually occur in the market basket data

Comparing the two approaches

◆ Approach 1: Triangular Matrix

➤ n = total number items

➤ Complexity: $O(n^2)$

➤ If

➤ If

➤ T

➤ T

◆ Approach 2

(but

➤ B


possible pairs actually occur

**Problem is if we have
too many items so the
pairs
do not fit into memory.
Can we do better?**

A-Priori Algorithm

A-Priori Algorithm – (1)

- ◆ A **two-pass** approach called *A-Priori* limits the need for main memory
- ◆ **Key idea:** *monotonicity*
 - If a set of items I appears at least s times, so does every **subset** J of I
- ◆ **Contrapositive for pairs:**
If item i does not appear in s baskets, then no pair including i can appear in s baskets
- ◆ **So, how does A-Priori find freq. pairs?**



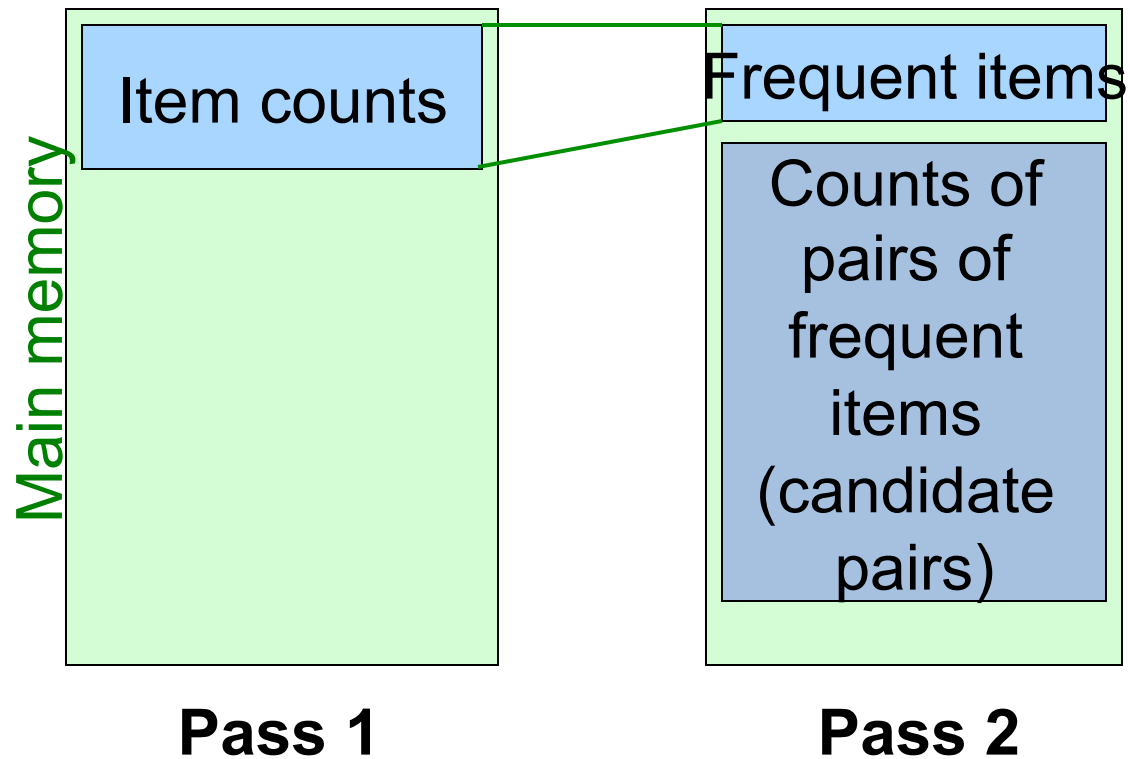
A-Priori Algorithm

- ◆ **Pass 1: Read baskets and count in main memory the occurrences of each item**
 - Requires only memory proportional to #items
- ◆ **Items that appear at least s times are the *frequent items***
 - At the end of pass 1, after the complete input file has been processed, check the count for each item
 - If $\text{count} > s$, then that item is frequent: saved for the next pass
- ◆ **Pass 1 identifies frequent itemsets (support $> s$) of size 1**

A-Priori Algorithm

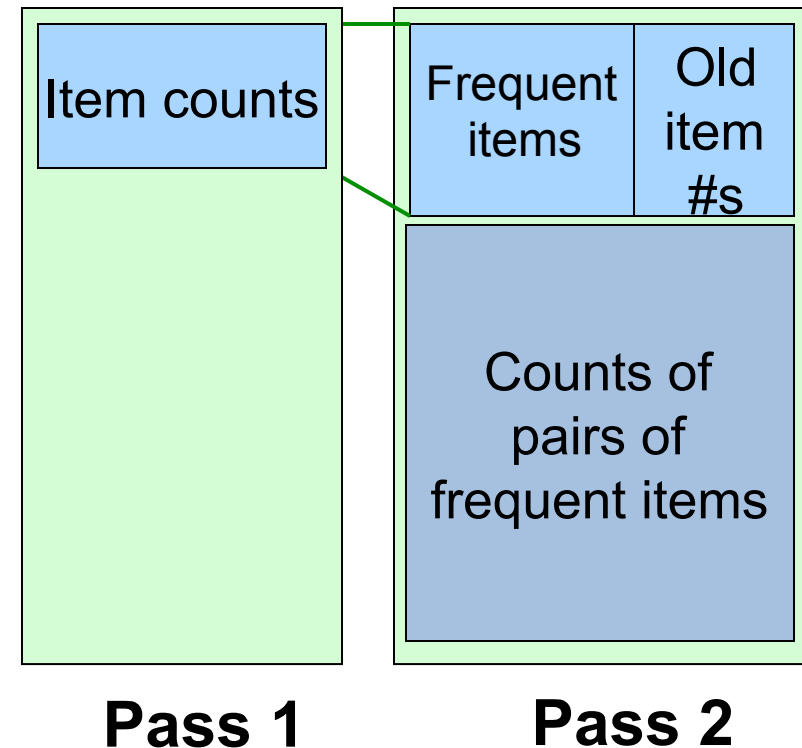
- ◆ **Pass 2: Read baskets again and count in main memory only those pairs of items where both were found in Pass 1 to be frequent**
- **Requires:**
 - **Memory proportional to square of *frequent* items only** (to hold counts of pairs)
 - **List of the frequent items from the first pass** (so you know what must be counted)
- ◆ **Pairs of items that appear at least s times are the *frequent pairs of size 2***
 - **At the end of pass 2, check the count for each pair**
 - **If $\text{count} > s$, then that pair is frequent**
- ◆ **Pass 2 identifies frequent pairs: itemsets of size 2**

Main-Memory: Picture of A-Priori



Detail for A-Priori

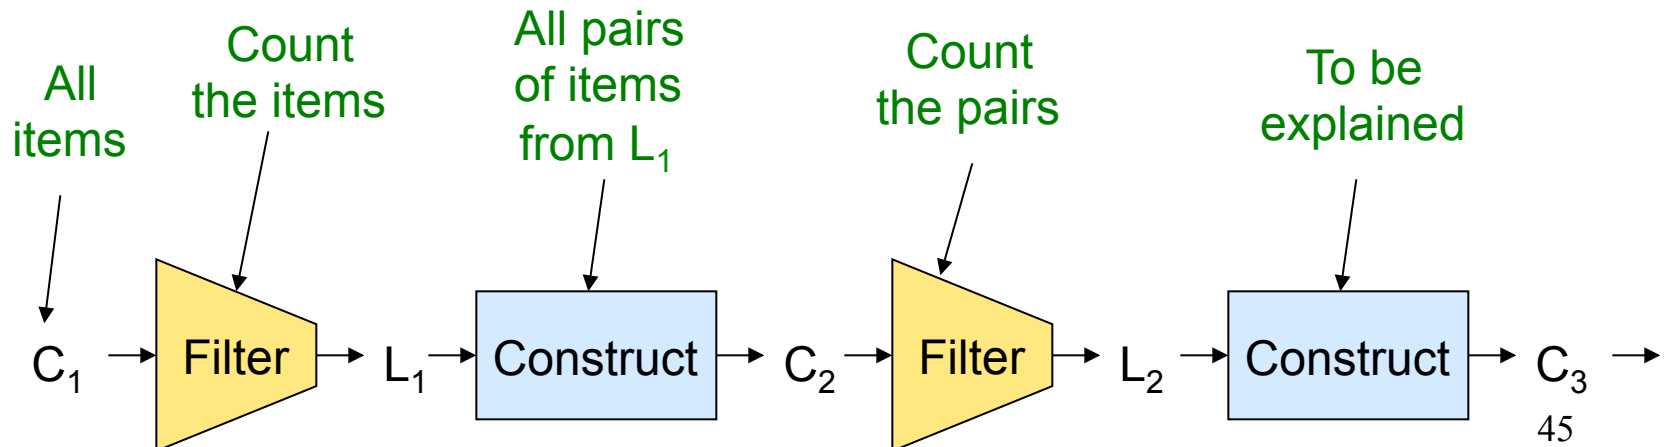
- ◆ You can use the triangular matrix method with n = number of frequent items
 - May save space compared with storing triples
- ◆ **Trick: re-number frequent items 1,2,...** and keep a table relating new numbers to original item numbers



What About Larger Frequent Itemsets?

Frequent Triples, Etc.

- ◆ For each k , we construct two sets of k -tuples (sets of size k):
 - C_k = *candidate k -tuples* = those that *might be frequent* sets (support $\geq s$) *based on information from the pass for $k-1$*
 - L_k = the set of *truly frequent k -tuples*



Recall: Example

$$B_1 = \{m, c, b\}$$

$$B_3 = \{m, c, b, n\}$$

$$B_5 = \{m, p, b\}$$

$$B_7 = \{c, b, j\}$$

$$B_2 = \{m, p, j\}$$

$$B_4 = \{c, j\}$$

$$B_6 = \{m, c, b, j\}$$

$$B_8 = \{b, c\}$$


◆ Frequent itemsets (s=3):

- $\{b\}, \{c\}, \{j\}, \{m\}$
- $\{b, m\} \quad \{b, c\} \quad \{c, m\} \quad \{c, j\}$
- $\{m, c, b\}$

Example

◆ Hypothetical steps of the A-Priori algorithm

- $C_1 = \{ \{b\} \{c\} \{j\} \{m\} \{n\} \{p\} \}$: all candidate items
- Count the support of itemsets in C_1
- Prune non-frequent: $L_1 = \{ b, c, j, m \}$
- Generate $C_2 = \{ \{b,c\} \{b,j\} \{b,m\} \{c,j\} \{c,m\} \{j,m\} \}$
- Count the support of itemsets in C_2
- Prune non-frequent: $L_2 = \{ \{b,m\} \{b,c\} \{c,m\} \{c,j\} \}$
- Generate $C_3 = \{ \{b,c,m\} \}$
- Count the support of itemsets in C_3
- Prune non-frequent: $L_3 = \{ \{b,c,m\} \}$



A-Priori for All Frequent Itemsets

- ◆ **One pass for each k (itemset size)**
- ◆ Needs room in main memory to count each candidate k -tuple
- ◆ For typical market-basket data and reasonable support (e.g., 1%), $k = 2$ requires the most memory